

AI-safety hackathon

Make careful AI-safety
communication the easy
default.

comm-ai-safety turns technical sources into reviewable
public communication—without losing the conditions that
make a claim true.

The failure happens when evidence crosses audiences.

Technical result

A result is bounded by model, task, prompting, baseline, uncertainty, and caveats.

Public artifact

A release, post, visual, or demo must stay clear, useful, and faithful under compression.

Avoidable distortion

Capability can become propensity; a benchmark can become real-world harm; a demo can become a misleading proxy.

source

translation

reader belief

A small workflow creates a practical quality-control layer.

1. Anchor claims

Build a compact ledger tied to primary sources, conditions, numbers, scope, and caveats.

2. Adapt responsibly

A researcher validates the main claims; the authoring skill fits them to one defined audience.

3. Audit before release

Check the finished artifact for fidelity, missing context, unsafe demos, rights, and accessibility.

evidence

communication

review

Three skills cover the channels where nuance most often disappears.

comm-public

Explainers, FAQs, social copy, visual briefs, and safer demo briefs for a named audience.
Built from primary sources and a named audience brief.



comm-press

Source-verifiable releases, media briefs, pitches, and journalist FAQs without manufactured hype.



comm-audit

A claim-by-claim check of conditions, caveats, quantitative context, safety, rights, and access.



One enforceable rule

No artifact is called audited or publishable until the required human checkpoints and material audit items are resolved.

Ready to use

Install the skills. Start with the
worked demo. Make a public
claim easier to verify.

Scan for the GitHub repo
Worked demo included
Human-reviewed claims

