AI-safety hackathon

Make careful AI-safety communication the easy default.

comm-ai-safety turns technical sources into reviewable public communication—without losing the conditions that make a claim true.

The failure happens when evidence crosses audiences.

Technical result
A result is bounded by model, task, prompting, baseline, uncertainty, and caveats.

Public artifact
A release, post, visual, or demo must stay clear, useful, and faithful under compression.

Avoidable distortion
Capability can become propensity; a benchmark can become real-world harm; a demo can become a misleading proxy.

source

translation

reader belief

A small workflow creates a practical quality-control layer.

1. Anchor claims
Build a compact ledger tied to primary sources, conditions, numbers, scope, and caveats.

2. Adapt responsibly
A researcher validates the main claims; the authoring skill fits them to one defined audience.

3. Audit before release
Check the finished artifact for fidelity, missing context, unsafe demos, rights, and accessibility.

evidence

communication

review

Three skills cover the channels where nuance most often disappears.

comm-public
Explainers, FAQs, social copy, visual briefs, and safer demo briefs for a named audience.
Built from primary sources and a named audience brief.

comm-press
Source-verifiable releases, media briefs, pitches, and journalist FAQs without manufactured hype.

comm-audit
A claim-by-claim check of conditions, caveats, quantitative context, safety, rights, and access.

One enforceable rule
No artifact is called audited or publishable until the required human checkpoints and material audit items are resolved.

Ready to use

Install the skills. Start with the worked demo. Make a public claim easier to verify.

github.com/chiffonng/comm-ai-safety
Worked case: Emergent Misalignment
Open Agent Skills · source-linked · human-in-the-loop